Siddh Mehta

AI / ML Engineer · Bengaluru, India

mehtasiddh011@gmail.com · siddhm11.github.io · linkedin.com/in/siddhm11 · github.com/siddhm11

SUMMARY

AI/ML engineer who builds retrieval systems, LLM tools and production ML pipelines. Built and run ResearchIT, hybrid search and multi-interest recommendations over 1.85M arXiv papers. About two years across Airbus, Samsung R&D (PRISM), Fynd and Acheron.

EXPERIENCE

Airbus India — ML / Data Engineer Feb 2026 – Present

Bengaluru · Intern Feb–Aug 2026; full-time since Aug 2026

- Engineered PySpark ingestion and reconciliation pipelines processing 43M+ rows in Palantir Foundry, with lineage, automated health checks and validation gates feeding downstream models and reporting.
- Developed milestone-delay and mean-time-to-repair models (LightGBM, MiniLM-L12-v2 embeddings, SQL features, per-class thresholds) with scheduled retraining.
- Built contract-analytics and consumption-forecasting pipelines for financial provisioning, and a self-service application with dashboards for global operations and finance teams.

Samsung R&D Institute India — PRISM research program Dec 2025 – Sep 2026

Remote

- Designed and benchmarked a two-stage document-image classifier (eight layout-aware classes) for constrained on-device inference, comparing CNN and transformer models on accuracy, latency and model size.

Fynd — AI/ML Engineering Intern May 2025 – Jul 2025

Mumbai

- Built a modular Python/LangGraph multi-agent framework with reusable abstractions, cutting pipeline setup time 40% and inference latency 60%.
- Implemented hybrid dense–sparse RAG with BGE-M3; served inference through REST APIs with unit tests and A/B evaluation harnesses, improving on baselines by 8–15% across three datasets.
- Fine-tuned a 1.5B-parameter Llama model with LoRA/QLoRA (PyTorch, PEFT), improving domain accuracy 9%.

Acheron Software — Software Engineering Intern · earlier Data Science Intern Jun 2024 – Apr 2025

Hyderabad / Remote · Software engineering intern Oct 2024 – Apr 2025; data science intern Jun – Aug 2024

- Shipped a LangChain Text-to-SQL application as an internal REST service (89% query accuracy) and a YOLOv8 asset-tagging pipeline with embedding similarity search (92% accuracy).
- Automated monitored data-processing and anomaly-detection workflows, saving about 20 minutes per upload cycle.

PROJECTS

ResearchIT — Personalized search and recommendations over 1.85M arXiv papers 2025 – Present

Solo · siddhm11-researchit.hf.space · source

- Hybrid search over 1.85M papers: BGE-M3 dense retrieval across three Qdrant shards plus SQLite FTS5, reciprocal rank fusion and a cross-encoder rerank; FastAPI/HTMX on Hugging Face Spaces with SQLite and Turso.
- Multi-interest feed: Ward clustering of saved papers into up to seven interests, importance-weighted quotas and per-interest MMR; traced and fixed a defect where one interest took 39 of 40 slots.
- Re-measured the reranker on production text before shipping it (MRR@10 +7.7% on 847 queries, 95% CI reported); restructured the FTS5 query to cut median lexical search from 877 to 288 ms.

Prompt Memory — Chrome extension for prompt rewriting with approved memory 2026

Co-built · source · evaluation corpus

- In-chat prompt rewriting on ChatGPT, Claude, Gemini, Perplexity and Grok (Manifest V3, FastAPI, Qdrant, MongoDB); context memory is written only after the user applies a rewrite.
- Built an 82-case adversarial evaluation corpus (21 failure categories) and revised the rewrite policy for source fidelity after it caught invented deadlines and limits.

SKILLS

Retrieval & LLMs: BGE-M3, Qdrant, SQLite FTS5, reciprocal rank fusion, cross-encoder reranking, RAG, LangGraph, LangChain, LoRA/QLoRA, evaluation design (MRR, nDCG)

ML & data: PyTorch, Transformers, scikit-learn, LightGBM, sentence-transformers, PySpark, SQL, Palantir Foundry

Engineering: Python, FastAPI, HTMX, Docker, SQLite/Turso, PostgreSQL, MongoDB, GitHub Actions, AWS, GCP

EDUCATION

Vellore Institute of Technology — B.Tech, CSE (Data Science) · CGPA 9.07 / 10 May 2026

AWS Certified Cloud Practitioner · Machine Learning Specialization · IBM Data Science Professional Certificate